

DAVID CLARK

Principal AI Platform Engineer

(650) 530-0517 | david.clark67@outlook.com | Rochester, MN, US
linkedin.com/in/david-clark-bab9053b9

SUMMARY

Principal AI Platform Engineer with 10 years of experience building production-grade AI systems and scalable full-stack platforms. Strong background in real-time data pipelines, machine learning deployment, and distributed systems, with a track record of increasing prediction accuracy by 28% and cutting latency by 35% in high-volume environments. Experienced in delivering end-to-end AI solutions, from feature engineering and NLP modeling to API integration and cloud deployment on AWS and Kubernetes. Comfortable working across ML, backend, and product layers, supporting systems that serve 100K+ daily users while driving measurable business impact.

TECHNICAL SKILLS

- **AI / ML Core Stack:** RAG Systems, Large Language Models (LLMs), Real-Time AI Systems, PyTorch, TensorFlow, NLP, Model Deployment, Feature Engineering, Deep Learning, Distributed AI Systems
- **Languages:** Python, TypeScript, JavaScript, Go, Java
- **Data & Streaming:** Kafka, Spark, Airflow, Distributed Data Pipelines
- **Cloud & Infra:** AWS (EC2, S3, Lambda, RDS, ECS), GCP, Kubernetes, Docker, Terraform, CI/CD pipelines
- **Databases:** Vector Databases, PostgreSQL, MongoDB, Redis, MySQL, Elasticsearch
- **Testing & Tools:** PyTest, Jest, Mocha, Cypress, TDD, Git, GitHub, GitLab, Jira, Prometheus, Grafana, Nginx

PROFESSIONAL EXPERIENCE

PredictHQ

Sep 2022 – Present

Principal AI Platform Engineer

- Built AI-driven event intelligence pipelines processing 10M+ global events, boosting forecasting accuracy by 28%.
- Trained machine learning models with XGBoost and PyTorch, increasing inference throughput by 22% in production.
- Cut data latency by 35% through real-time ingestion pipelines using Kafka and AWS.
- Elevated feature quality by 18% via advanced temporal and geospatial feature engineering.
- Established automated model retraining and monitoring, lowering model drift incidents by 40%.
- Built NLP-based classification systems, raising event tagging accuracy by 25%.
- Streamlined hybrid batch and streaming inference, trimming infrastructure costs by 20% while maintaining performance.
- Supported AI-powered APIs handling 100K+ daily requests with stable low-latency performance.
- Accelerated API response times by 30% through efficient service architecture and caching strategies.
- Deployed CI/CD pipelines for ML services, shortening deployment time by 50%.
- Built observability systems with Prometheus and Grafana, sustaining 99.9% uptime.

- Mentored engineers on ML system design, strengthening delivery speed and production reliability.

Prezi

Mar 2020 – Aug 2022

AI Platform Engineer

- Boosted user engagement by 20% by creating recommendation models for personalized content.
- Built NLP pipelines for slide content analysis, increasing relevance scoring accuracy by 25%.
- Built user behavior prediction models that raised retention metrics by 15%.
- Accelerated experimentation cycles by 30% through A/B testing frameworks for AI features.
- Streamlined data pipelines with Spark, reducing processing time by 35%.
- Productionized ML models with cross-functional teams, cutting deployment friction by 40%.
- Enabled real-time content suggestions using React and Flask-based AI features.
- Delivered backend APIs for ML services, supporting scalable integration across applications.

Prezi

Feb 2018 – Mar 2020

Senior Full Stack Engineer

- Boosted frontend performance by 30% by building scalable React and Redux applications.
- Developed backend services supporting 50K+ daily active users using Node.js and Python.
- Increased query performance by 25% through efficient PostgreSQL schema design.
- Cut page load time by 40% with frontend optimization and lazy loading strategies.
- Strengthened system reliability by transitioning from monolithic to microservices architecture.

Trbhi

June 2016 – Dec 2017

Full Stack Engineer

- Built 5+ scalable web applications, increasing transaction efficiency.
- Enhanced frontend responsiveness using React and Node.js, reducing load times.
- Re-architected database systems and migrated legacy infrastructure to MongoDB for improved scalability.
- Improved backend services and implemented CI/CD pipelines with Jenkins, accelerating release cycles.
- Introduced unit and integration testing across 8+ services, improving system stability.
- Integrated third-party services, expanding platform capabilities while maintaining 99.9% availability.

EDUCATION

Master of Science in Computer Science

Sep 2014 – June 2016

Stratford University

Bachelor of Science in Computer Science

Jan 2011 – June 2014

Stratford University